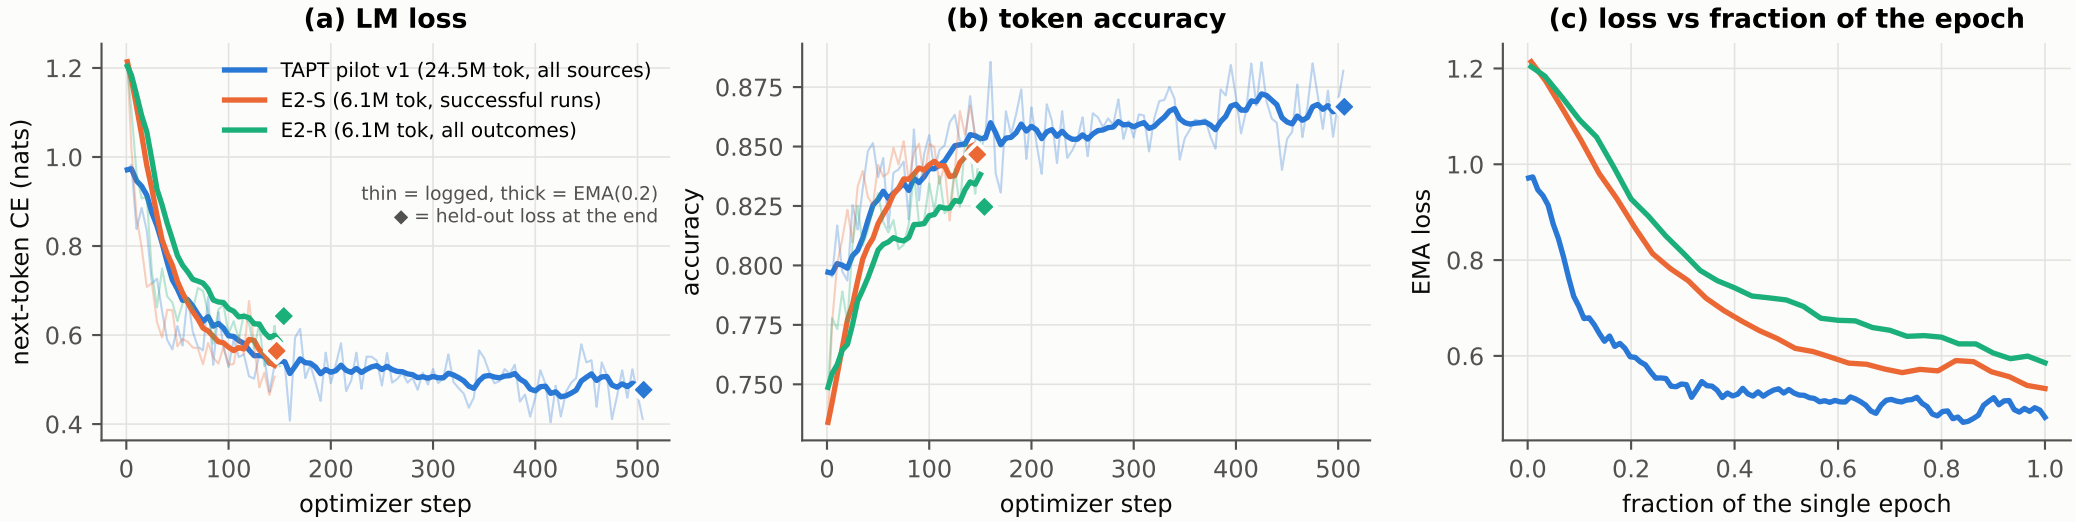
Text mid-training of Qwen3-8B (MS-Swift continued pre-training, LoRA r64 all-linear, 1 epoch)
(a) LM loss
TAPT pilot v1 (24.5M tok, all sources)
E2-S (6.1M tok, successful runs)
E2-R (6.1M tok, all outcomes)
thin = logged, thick = EMA(0.2)
◆ = held-out loss at the end
next-token CE (nats)
optimizer step
(b) token accuracy
accuracy
optimizer step
(c) loss vs fraction of the epoch
EMA loss
fraction of the single epoch