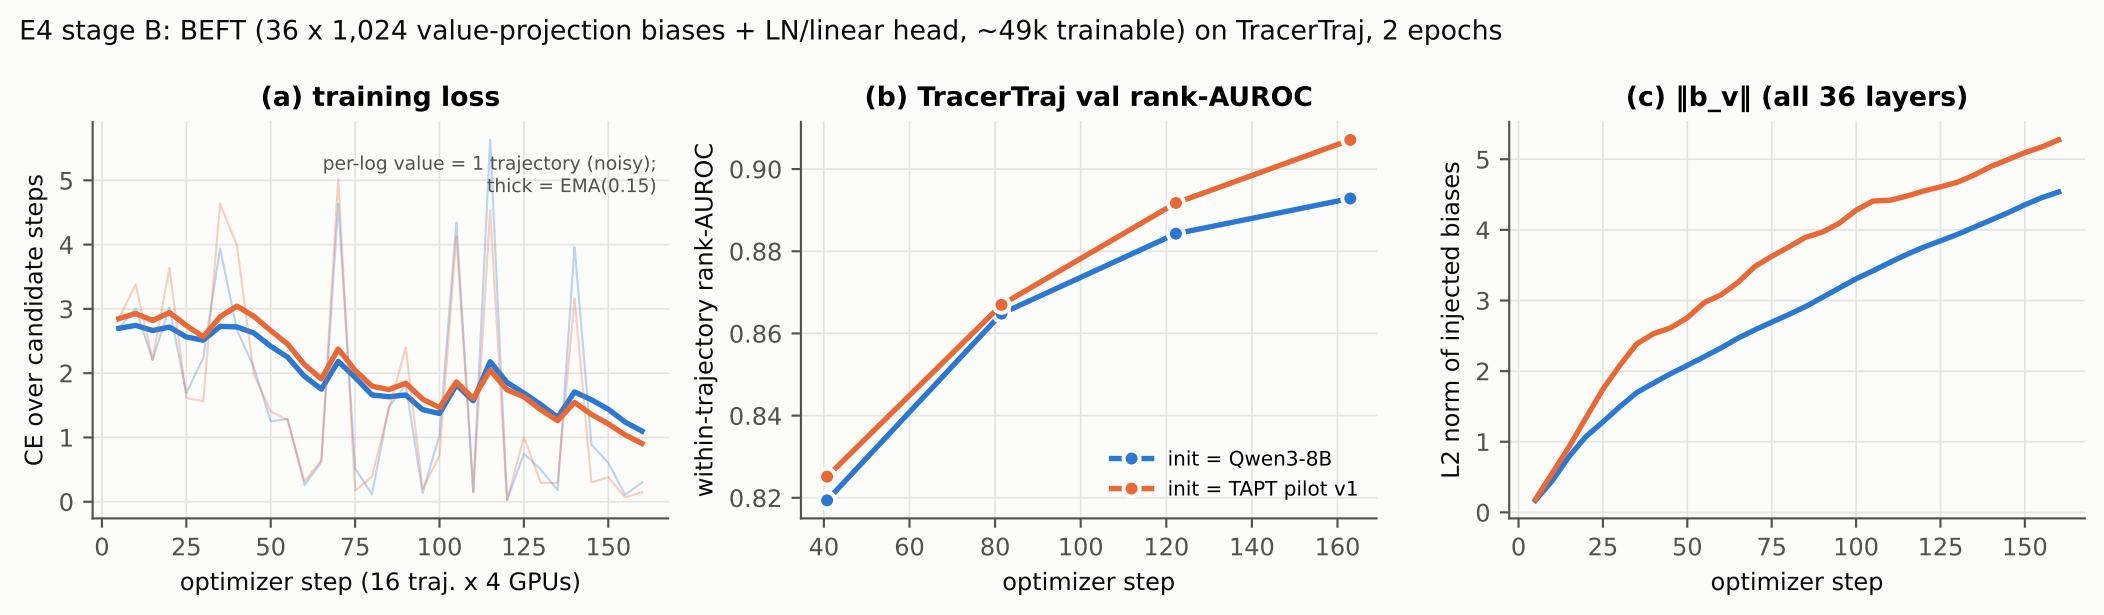
E4 stage B: BEFT (36 x 1,024 value-projection biases + LN/linear head, ~49k trainable) on TracerTraj, 2 epochs
(a) training loss
per-log value = 1 trajectory (noisy);
thick = EMA(0.15)
CE over candidate steps
optimizer step (16 traj. x 4 GPUs)
(b) TracerTraj val rank-AUROC
within-trajectory rank-AUROC
optimizer step
init = Qwen3-8B
init = TAPT pilot v1
(c) ‖b_v‖ (all 36 layers)
L2 norm of injected biases
optimizer step